

CODING LOOPS WITH "apply" - LECTURE

DSC 205 - Advanced introduction to data science

Marcus Birkenkrahe

February 24, 2023

README



Figure 1: Photo by Brett Jordan on Unsplash

- Download the **raw** file to practice during the lecture from GitHub, save it as `7_loop_apply_practice.org` and upload it to Canvas later.
- To test your Emacs mettle, open it on the CMD line with the command `emacs -nw` (no graphics - not needed for this exercise).

`apply`, `lapply`, `sapply`, `tapply`, `vapply`, `mapply`



Figure 2: Family by Rajiv Perera on Unsplash

- The **apply** family of functions allows implicit looping over subsets of vectors, matrices or arrays to apply a function a subset of their elements.
- The different flavors of **apply** are:
 1. **apply** applies a function to a dataset's margin (segment)
 2. **tapply** to apply a function to subsets defined in terms of **factor** vectors, i.e. sliced by categorical variable values.
 3. **lapply** operates member by member on a **list**.
 4. **sapply** to return simplified **lapply** results.
 5. **vapply** if you know the data type you're expecting as a return.
 6. **mapply** as a multi-variate version of **sapply**

- All **apply** type functions simplify coding enormously by replacing **for** loop constructions.
- All **apply** type functions allow for additional arguments (...) to be passed to FUN.

apply - implicit looping over arrays

- The **apply** function returns a vector or array or list of values by applying a function to the **MARGIN** of an array or matrix.

```
apply( X = data # array with positive dim(X)
      MARGIN = subset, # aka row/column/layer etc.
      FUN = function) # includes your own function
```

- Does **apply** have any other arguments? Find out!

```
args(apply)
```

```
function (X, MARGIN, FUN, ..., simplify = TRUE)
NULL
```

What's an array?

- scalars: 0-dim arrays, e.g. the number 0
- vectors: 1-dim arrays, e.g. `c("a","vector")`
- matrices: 2-dim arrays, e.g. `matrix(1:9,3,3)`
- arrays: n-dim, e.g. `array(1:9,dim=c(3,3,2))`

```
array(1:18,dim=c(3,3,2))
```

```
, , 1
```

```
      [,1] [,2] [,3]
[1,]    1    4    7
[2,]    2    5    8
[3,]    3    6    9
```




Figure 3: R data structures (source: Ceballos, 2013)

, , 2

	[,1]	[,2]	[,3]
[1,]	10	13	16
[2,]	11	14	17
[3,]	12	15	18

- The **MARGIN** index follows the positional order of the dimension for matrices and arrays:

MARGIN	DATA STRUCTURE
1	rows
2	columns
3	layers
4	blocks

apply example: matrix

- Create a 4 x 3 matrix with the elements 1:12:

```
foo <- matrix(1:12,4,3) # default byrow=FALSE
foo
```

	[,1]	[,2]	[,3]
[1,]	1	5	9
[2,]	2	6	10
[3,]	3	7	11
[4,]	4	8	12

- Find the sum of each row of `foo`. What about `sum(foo)`?

```
sum(foo) # this sums up ALL elements
```

```
[1] 78
```

- Loop over the rows of `foo`:

```
row.totals <- rep(NA, times=nrow(foo)) # initialize counter
for (i in 1:nrow(foo)) {
  row.totals[i] <- sum(foo[i,]) # sum over i-th row
}
row.totals
```

```
[1] 15 18 21 24
```

- Much shorter with `apply`:

```
apply(
  X = foo,
  MARGIN = 1, # MARGIN = 1 (rows), 2 (cols), 3 (layers), 4 (blocks)
  FUN = sum)
```

```
[1] 15 18 21 24
```

- To sum over columns instead, change `MARGIN` to 2.

```
apply(
  X = foo,
  MARGIN = 2,
  FUN = sum)
```

```
[1] 10 26 42
```

- You can pass additional arguments to any **apply** function: e.g. you can use the function **sort** and specify it to be **decreasing**:

```
apply(
  X = foo,
  MARGIN = 1,
  FUN = sort,
  decreasing = TRUE)
```

	[,1]	[,2]	[,3]	[,4]
[1,]	9	10	11	12
[2,]	5	6	7	8
[3,]	1	2	3	4

apply example: array

- Create a 3 x 2 x 2 array **bar** with the elements 1:18

```
bar <- array(1:18, dim=c(3,3,2))
bar
```

```
, , 1
```

	[,1]	[,2]	[,3]
[1,]	1	4	7
[2,]	2	5	8
[3,]	3	6	9

```
, , 2
```

	[,1]	[,2]	[,3]
[1,]	10	13	16
[2,]	11	14	17
[3,]	12	15	18

- Put differently, **bar** has 2 layers of 3 x 3 matrices. What does the following call do?¹

¹The **apply** call extracts the diagonal elements for each of the 2 layers with **diag**. Each call to **diag** of a matrix returns a vector and these vectors are returned as columns of a new matrix.

```
baz <- apply(bar,3,FUN=diag)
baz
```

```
      [,1] [,2]
[1,]     1    10
[2,]     5    14
[3,]     9    18
```

- Check the dimensions and class of baz:

```
dim(baz)
class(baz)
is.matrix(baz)
is.array(baz)

[1] 3 2
[1] "matrix" "array"
[1] TRUE
[1] TRUE
```

tapply - slicing data by categories

- `tapply` performs operations on subsets defined by **factor** vectors
- Simple example: compute the **mean** tooth length by supply category in the `ToothGrowth` dataset:

```
tapply(X = ToothGrowth$len, # length of guinea pig teeth
      INDEX = ToothGrowth$supp, # OJ or VC supply
      FUN = mean) # arithmetic average
```

```
      OJ      VC
20.66333 16.96333
```

- The result returns the average length for guinea pigs supplied with orange juice (OJ) and vitamin C (VC).
- Here's another example (data source: Kaggle)²:

²Astonishingly, some websites are trying to sell these (freely available) data for US\$100.00 (see [here](#)).

1. read web data on diamond pricing (with strings as factors)
2. display structure of data table
3. display first five records

```
dia.url <- "https://raw.githubusercontent.com/birkenkrahe/ds2/main/data/diamonds.csv"
diamonds <- read.csv(dia.url, stringsAsFactors=TRUE)
str(diamonds)
head(diamonds)
```

```
'data.frame': 53943 obs. of 11 variables:
 $ X      : int  1 2 3 4 5 6 7 8 9 10 ...
 $ carat  : num  0.23 0.21 0.23 0.29 0.31 0.24 0.24 0.26 0.22 0.23 ...
 $ cut    : Factor w/ 5 levels "Fair","Good",...: 3 4 2 4 2 5 5 5 1 5 ...
 $ color  : Factor w/ 7 levels "D","E","F","G",...: 2 2 2 6 7 7 6 5 2 5 ...
 $ clarity: Factor w/ 8 levels "I1","IF","SI1",...: 4 3 5 6 4 8 7 3 6 5 ...
 $ depth  : num  61.5 59.8 56.9 62.4 63.3 62.8 62.3 61.9 65.1 59.4 ...
 $ table  : num  55 61 65 58 58 57 57 55 61 61 ...
 $ price  : int  326 326 327 334 335 336 336 337 337 338 ...
 $ x      : num  3.95 3.89 4.05 4.2 4.34 3.94 3.95 4.07 3.87 4 ...
 $ y      : num  3.98 3.84 4.07 4.23 4.35 3.96 3.98 4.11 3.78 4.05 ...
 $ z      : num  2.43 2.31 2.31 2.63 2.75 2.48 2.47 2.53 2.49 2.39 ...

  X carat      cut color clarity depth table price    x    y    z
1 1  0.23    Ideal     E    SI2   61.5    55   326 3.95 3.98 2.43
2 2  0.21  Premium     E    SI1   59.8    61   326 3.89 3.84 2.31
3 3  0.23     Good     E    VS1   56.9    65   327 4.05 4.07 2.31
4 4  0.29  Premium     I    VS2   62.4    58   334 4.20 4.23 2.63
5 5  0.31     Good     J    SI2   63.3    58   335 4.34 4.35 2.75
6 6  0.24 Very Good     J   VVS2   62.8    57   336 3.94 3.96 2.48
```

- Using `tapply`, you can add up the total value of the diamonds for the full data set but separated according to `color` (key coded):

```
tapply(
  X = diamonds$price,
  INDEX = diamonds$color,
  FUN = sum)

      D      E      F      G      H      I      J
21476439 30148457 35545622 45158240 37257301 27608146 14949281
```


`lapply` - cycling through lists

- `lapply` operates member by member on a `list` and returns a `list`:

```
baz <- list(  
  aa = c(3.4,1),  
  bb = matrix(1:4,2,2),  
  cc = matrix(c(T,T,F,T,F,F),3,2),  
  dd = "string here",  
  ee = matrix(c("red","green","blue","yellow")))
```

- Check for matrices in the list `baz`:

```
baz <- list(  
  aa = c(3.4,1),  
  bb = matrix(1:4,2,2),  
  cc = matrix(c(T,T,F,T,F,F),3,2),  
  dd = "string here",  
  ee = matrix(c("red","green","blue","yellow")))  
lapply(  
  X = baz,  
  FUN = is.matrix)
```

```
$aa  
[1] FALSE
```

```
$bb  
[1] TRUE
```

```
$cc  
[1] TRUE
```

```
$dd  
[1] FALSE
```

```
$ee  
[1] TRUE
```

- No margin or index information is required. R knows how to apply `FUN` to each member of the list, and returns a `list`. Fun!

sapply - simplified cycling

- `sapply` (`s = "simplified"`) returns the same results as `lapply` but in an array form:

```
baz <- list(
  aa = c(3.4,1),
  bb = matrix(1:4,2,2),
  cc = matrix(c(T,T,F,T,F,F),3,2),
  dd = "string here",
  ee = matrix(c("red","green","blue","yellow")))
sap <- sapply(
  X = baz,
  FUN = is.matrix)

sap
is.vector(sap)    # sap is a named vector

      aa      bb      cc      dd      ee
FALSE  TRUE  TRUE FALSE  TRUE
[1] TRUE
```

- `baz` has a `names` attribute that is copied to the corresponding entries of the returned object:

```
attributes(sap)
names(sap)
str(sap)

$names
[1] "aa" "bb" "cc" "dd" "ee"
[1] "aa" "bb" "cc" "dd" "ee"
  Named logi [1:5] FALSE TRUE TRUE FALSE TRUE
- attr(*, "names")= chr [1:5] "aa" "bb" "cc" "dd" ...
```

- If we did not have `sapply`, you could `unlist` the result of `lapply`:

```
unlist(lapply(baz,is.matrix))
sapply(baz,is.matrix)

      aa      bb      cc      dd      ee
FALSE  TRUE  TRUE FALSE  TRUE
      aa      bb      cc      dd      ee
FALSE  TRUE  TRUE FALSE  TRUE
```

SOMEDAY `vapply` - simplified cycling with safety check

- Read the help file and this tutorial (Treadway, 2020).

SOMEDAY `mapply` - multivariate version of `sapply`

- Read the help file and this tutorial (Zach, 2021).

TODO Exercises



1. Write an implicit loop that calculates the product of all the column elements of the matrix returned by the call to `apply(foo, 1, sort, decreasing=TRUE)` where `foo` is `matrix(1:12,4.3)`.
Tip: To multiply numbers, you can use the function `prod`.
2. Convert the following `for` loop to an implicit loop that does exactly the same thing. Here, `t` transposes its matrix argument.

Bonus: compare the results of the two operations without looking.

```
matlist <- list(
  matrix(c(T,F,T,T),2,2),
  matrix(c("a","c","b","z","p","q"),3,2),
  matrix(1:8,2,4))
matlist # original list

for (i in 1:length(matlist)) {
  matlist[[i]] <- t(matlist[[i]])
}
matlist # transposed list

[[1]]
      [,1] [,2]
[1,]  TRUE TRUE
[2,] FALSE TRUE

[[2]]
      [,1] [,2]
[1,] "a"  "z"
[2,] "c"  "p"
[3,] "b"  "q"

[[3]]
      [,1] [,2] [,3] [,4]
[1,]    1    3    5    7
[2,]    2    4    6    8
[[1]]
      [,1] [,2]
[1,] TRUE FALSE
[2,] TRUE  TRUE

[[2]]
      [,1] [,2] [,3]
[1,] "a"  "c"  "b"
[2,] "z"  "p"  "q"

[[3]]
      [,1] [,2]
```

[1,]	1	2
[2,]	3	4
[3,]	5	6
[4,]	7	8

Glossary

TERM	MEANING
<code>apply</code>	apply function to the margin of a dataset X
<code>tapply</code>	apply function to subsets grouped by factor
<code>lapply</code>	apply function to list members, return list
<code>sapply</code>	simplified <code>lapply</code> , returns vector
<code>vapply</code>	apply when you know the return datatype
<code>mapply</code>	multivariate version of <code>sapply</code>

References

- Ceballos, M. (2013). Data structure. URL: venus.ifca.unican.es.
- Davies, T.D. (2016). The Book of R. NoStarch Press.
- Treadway, A. (20 Oct 2020). Why you should use `vapply` in R. URL: theautomatic.net.
- Zach (Dec 7, 2021). How to Use the `mapply()` Function in R (With Examples). URL: statology.org.